

Introduction to **make**, BLAS and LAPACK

李阳 周嘉鑫

March 19, 2020

Contents

1 **make**

2 BLAS and LAPACK

Overview of **make**

- **make** is a program for directing recompilation, which automatically determines which pieces of a large program need to be recompiled, and issues commands to recompile them.
- Why **make**?
 - ① Reduce unnecessary repeated work.
 - ② **make** provides a way of efficiently codifying and organizing the instructions to be executed when building a program.
- You need a file called **Makefile** to tell **make** what to do. Most often, the makefile tells **make** how to compile and link a program.
- Issue the following command to install **make**:

```
sudo apt-get install make
```

The Syntax of Makefiles

- A makefile consists of a set of dependencies and rules.

```
target : prerequisites
      rule
```

- *Note: put a Tab character at the beginning of the rule line!*
- A target is usually an executable or object file. It can also be the name of an action to carry out, such as 'clean'
- A prerequisite is a file that is used as input to create the target. A target often depends on several files.
- The *rules* describe how to create the target from the prerequisites.

A Simple Makefile

```
main : main.o add.o sub.o mul.o
    g++ -o main main.o add.o sub.o mul.o
main.o : main.cpp arithmetic.h
    g++ -c main.cpp
add.o : add.cpp arithmetic.h
    g++ -c add.cpp
sub.o : sub.cpp arithmetic.h
    g++ -c sub.cpp
mul.o : mul.cpp arithmetic.h
    g++ -c mul.cpp
```

Variables in Makefile

- You define a variable in a makefile by writing `VARIABLE_NAME=value`, then accessing the value of `VARIABLE_NAME` by writing either `$(VARIABLE_NAME)` or `${VARIABLE_NAME}`.
- A comment in a makefile starts with `#` and continues to the end of the line.
- Some examples:
`CC = g++ # Which compiler`
`# Options for development`
`CPPFLAGS = -g -std=c++11`
`# Options for linking dynamic libraries`
`LDFLAGS = -llapacke -lblas`
`LIBS = main.o add.o sub.o mul.o`

Automatic Variables

- `$@` The name of the current target
- `$<` The name of the first prerequisite
- `$$` The names of all the prerequisites, with spaces between them
- Some examples:

```
main : main.o add.o sub.o mul.o
      $(CC) -o $@ $(CPPFLAGS) $$
main.o : main.cpp arithmetic.h
        $(CC) $(CPPFLAGS) -c $<
```

Built-in Rules

- **make** has a large number of built-in rules that can significantly simplify makefiles.
- To check this, create a traditional Hello World program called `hello.cpp` and compile it using
`make hello`
- You can print the built-in rules by issuing
`make -p`

Putting it all together

```
CC = g++
CPPFLAGS = -g -std=c++11
LIBS = main.o add.o sub.o mul.o
RM = rm -rf

main : $(LIBS)
    $(CC) -o $@ $(CPPFLAGS) $(LIBS)

.PHONY : clean

clean :
    $(RM) *.o main
```

Makefile for Your Homework

```
default: homeworkTemplate.pdf
```

```
%.pdf : %.tex  
    pdflatex $<  
    pdflatex $<
```

```
clean :  
    $(RM) *.aux *.log
```

```
realclean :  
    $(MAKE) clean  
    $(RM) *.pdf
```

Some useful options to **make**

- **-k** Continue as much as possible after an error.
- **-n** Print what **make** would have done without actually doing it.
- **-s** Do not print the rules as they are executed.
- **-f filename** Read the file named `filename` as a makefile.

Libraries for Linear Algebra

- Linear algebra is a fundamental building block of nearly every scientific software package.
- Reasons for using libraries instead of coding an algorithm from scratch:
 - ① **programmer productivity**: focus on **what** instead of **how**;
 - ② **code quality**: a library is usually specifically tuned for your computer architecture.

Overview of BLAS and CBLAS

- The BLAS(Basic Linear Algebra Subprograms) library contains high quality building block routines for performing basic vector and matrix operations.
- BLAS routines are written in Fortran.
- CBLAS provides a C/C++ language interface to BLAS routines.
- Install the package `libblas-dev` from Synaptic Package Manager.

BLAS

The BLAS libraries are decomposed into problems at three levels:

- BLAS Level 1 operations work on vectors and have a typical complexity of n operations for a vector of length n . Typical examples are vector addition, computing an inner product, or applying Givens rotations.
- BLAS Level 2 operations are of the matrix-vector type, with n^2 complexity, such as the matrix-vector product.
- BLAS Level 3 operations are between two matrices, the typical example being the matrix-matrix product, which has $O(n^3)$ complexity.

Naming Conventions for BLAS Routines

$$XYZZZZ : \begin{cases} X : & \text{precision} \\ YY : & \text{matrix type} \\ ZZZ : & \text{operation} \end{cases}$$

- The “precision” character X can be S for single, D for double, C for single-precision complex, and Z for double-precision complex.
- The two characters YY reflects the matrix argument type. For instance, GE stands for general rectangular, GT for general triangular, SY for symmetric, TR for triangular, etc.
- The (up to three) characters ZZZ describe the operation abbreviated with shorthands, including MV for matrix-vector product or MM for matrix-matrix product.
- SGEMV: single precision, general matrix, matrix-vector product
- ZTRMM: double-precision complex, triangular matrix, matrix-matrix product

CBLAS

- In CBLAS, the Fortran routine names are prefixed with `cblas_`. Names of all CBLAS functions are in lowercase letters.
- `SGEMV` becomes `cblas_sgemv`, and `ZTRMM` becomes `cblas_ztrmm`.
- All CBLAS function declarations are in the file `cblas.h`.

BLAS Level 1 Example

Compute the inner product of two vectors:

```
double cblas_ddot(const int n, const double *x, const int
    incx, const double *y, const int incy)
```

`incx` and `incy` specify the increment for the elements of `x` and `y`, respectively.

Create a file named `ddot.cpp` with the following contents:

```
#include <iostream>
#include <blas.h>

int main(int argc, char *argv[]) {
    double x[3] = {1.1, 1.2, 1.3};
    double y[3] = {1.0, 2.0, 3.0};
    std::cout << "The inner product of x and y is ";
    std::cout << cblas_ddot(3, x, 1, y, 1) << std::endl;
}
```

Issue the following command:

```
g++ -o ddot ddot.cpp -lblas
./ddot
```

BLAS Level 2 Example

Compute a matrix-vector product using a general matrix: $\mathbf{y} \leftarrow \alpha A\mathbf{x} + \beta\mathbf{y}$

```
void cblas_dgemv(const CBLAS_ORDER order,  
                const CBLAS_TRANSPOSE trans,  
                const int m, const int n,  
                const double alpha, const double *A, const int lda,  
                const double *x, const int incx,  
                const double beta, double *y, const int incy)
```

CBLAS_ORDER

- There are two general methods of storing a two dimensional matrix in linear (one dimensional) memory: column-wise (column major order) or row-wise (row major order).
- Consider a general matrix $A = (a_{ij})_{1 \leq i \leq M, 1 \leq j \leq N}$. In column major order, the matrix elements are located in memory according to this sequence:

$$a_{1,1}, a_{2,1}, \dots, a_{M,1}, a_{1,2}, a_{2,2}, \dots, a_{M,2}, \dots, a_{1,N}, a_{2,N}, \dots, a_{M,N},$$

In row major order, the matrix elements are located in memory according to this sequence:

$$a_{1,1}, a_{1,2}, \dots, a_{1,N}, a_{2,1}, a_{2,2}, \dots, a_{2,N}, \dots, a_{M,1}, a_{M,2}, \dots, a_{M,N}$$

- `order=CblasRowMajor`: row-major. `order=CblasColMajor`: column-major.

CBLAS_TRANSPOSE

- `trans=CblasNoTrans`: $\mathbf{y} \leftarrow \alpha \mathbf{A}\mathbf{x} + \beta \mathbf{y}$.
- `trans=CblasTrans`: $\mathbf{y} \leftarrow \alpha \mathbf{A}^T \mathbf{x} + \beta \mathbf{y}$.
- `trans=CblasConjTrans`: $\mathbf{y} \leftarrow \alpha \mathbf{A}^H \mathbf{x} + \beta \mathbf{y}$.
- `trans=CblasConjNoTrans`: $\mathbf{y} \leftarrow \alpha \bar{\mathbf{A}} \mathbf{x} + \beta \mathbf{y}$.

Leading Dimension Parameter

- A leading dimension parameter allows use of BLAS routines on a submatrix of a larger matrix.
- Consider a submatrix B extracted from the original matrix A :

$$B = \begin{bmatrix} a_{i_0+1,j_0+1} & a_{i_0+1,j_0+2} & \cdots & a_{i_0+1,j_0+L} \\ a_{i_0+2,j_0+1} & a_{i_0+2,j_0+2} & \cdots & a_{i_0+2,j_0+L} \\ \vdots & \vdots & \ddots & \vdots \\ a_{i_0+K,j_0+1} & a_{i_0+K,j_0+2} & \cdots & a_{i_0+K,j_0+L} \end{bmatrix}$$

- To specify matrix B , BLAS routines require four parameters:
 - ① the number of rows K ;
 - ② the number of columns L ;
 - ③ a pointer to the start of the array containing elements of B ;
 - ④ the leading dimension of the array containing elements of B .


Leading Dimension Parameter

The leading dimension depends on the storage order of the matrix A :

- Column major order

Leading dimension is M , i.e., the number of rows of matrix A .


Starting address: offset by $i_0 + j_0 M$ from $a_{1,1}$.



- Row major order

Leading dimension is N , i.e., the number of columns of the matrix A .

Starting address: offset by $i_0 N + j_0$ from $a_{1,1}$.



```
#include <iostream>
#include <algorithm>
#include <iterator>
#include <cbblas.h>
using namespace std;

int main(int argc, char *argv[]) {
    double mat[4*3] = {1, 2, 0, 1,
                       2, 0, 1, 1,
                       0, 3, 2, 2};
    double x[3] = {1, 2, 3};
    double y[3] = {0, 0, 0};
    cbblas_dgemv(CblasColMajor, CblasNoTrans,
                 3, 3, 1.0, mat, 4, x, 1, 0.0, y, 1);
    copy(begin(y), end(y),
         ostream_iterator<double>(cout, " "));
    cout << endl;
}
```


BLAS Level 3 Example

Compute a matrix-matrix product with general matrices:

```
void cblas_dgemm(const CBLAS_ORDER order, const
CBLAS_TRANSPOSE transa, const CBLAS_TRANSPOSE transb,
const int m, const int n, const int k,
const double alpha, const double *a, const int lda,
const double *b, const int ldb, const double beta,
double *c, const int ldc)
```

Try it yourself!

For more information, refer to <http://netlib.org/blas/>.

Overview of LAPACK

- LAPACK(Linear Algebra PACKage) is written in Fortran 90 and provides routines for solving systems of linear equations, linear least-square problems, eigenvalue and singular value problems, and performing a number of related computational tasks.
- LAPACK routines are written so that as much as possible of the computation is performed by calls to the BLAS.
- Naming convention of LAPACK follows from that of BLAS.

Overview of LAPACK

- LAPACK provides a C/C++ language interface to LAPACK routines.
- In LAPACK, the Fortran routine names are prefixed with LAPACK_.
- All LAPACK function declarations are in the file `lapack.h`.
- Install the package `liblapack-dev` from Synaptic Package Manager.

LAPACKE_dgesv

```
lapack_int LAPACKE_dgesv(int matrix_layout,  
    lapack_int n, lapack_int nrhs,  
    double* A, lapack_int lda,  
    lapack_int* ipiv,  
    double* B, lapack_int ldb)
```

This routine solves for X the system of linear equations $AX = B$, where A is a square matrix, the columns of B are individual right-hand sides, and the columns of X are the corresponding solutions.

The LU decomposition with partial pivoting and row interchanges is used to factor A as $A = PLU$, where P is a permutation matrix, L is unit lower triangular, and U is upper triangular. The factored form of A is then used to solve the system of equations $AX = B$.

LAPACKE_dgesv

- `matrix_layout` specifies whether matrix storage is row major(`LAPACK_ROW_MAJOR`) or column major(`LAPACK_COL_MAJOR`).
- Side effects:
 - ① A is overwritten by the factors L and U from the factorization $A = PLU$; the unit diagonal elements of L are not stored.
 - ② B is overwritten by the solution matrix X .
 - ③ `ipiv` is the pivot vector that defines the permutation matrix.
- `LAPACKE_dgesv` returns a value `info`.
 - ① If `info=0`, the execution is successful.
 - ② If `info=-i`, parameter i had an illegal value.
 - ③ If `info=i`, $U_{i,i}$ is exactly zero. The factorization has been completed, but the factor U is exactly singular, so the solution could not be computed.

```
#include <iostream>
#include <lapacke.h>
#include <algorithm>
#include <iterator>
using namespace std;

int main(int argc, char *argv[]) {
    double A[9] = {2, -2, -1,
                  -2, 4, 8,
                  6, 3, 4};
    double b[3] = {16, 0, -1};
    int ipiv[3];
    LAPACKE_dgesv(LAPACK_COL_MAJOR, 3, 1, A, 3, ipiv, b, 3);
    cout << "The solution is : \n";
    copy(begin(b), end(b),
         ostream_iterator<double>(cout, " "));
    cout << endl;
}
```

LAPACKE_dsterf

Compute all eigenvalues of a real symmetric tridiagonal matrix.

```
lapack_int LAPACKE_dsterf(lapack_int n, double *d, double *e)
```

Input Parameters:

- n : The order of the matrix
- **d**: **d** contains the diagonal elements
- **e**: **e** contains the off-diagonal elements

Side Effects:

- **d**: The n eigenvalues in ascending order.

```
#include <iostream>
#include <lapacke.h>
#include <algorithm>
#include <iterator>
using namespace std;

int main(int argc, char *argv[]) {
    double d[4] = {2, 2, 2, 2};
    double e[3] = {-1, -1, -1};
    auto info = LAPACKE_dsterf(4, d, e);
    if(info==0) {
        copy(begin(d), end(d),
            ostream_iterator<double>(cout, " "));
    }
    else {
        cout << "LAPACKE_dsterf failed\n";
    }
}
```


References

- GNU Make Manual
- Beginning Linux Programming
- Introduction to Scientific and Technical Computing
- <http://netlib.org/blas/>
- <http://www.netlib.org/lapack/>